

# GNR638 Machine Learning for Remote Sensing II

## Mini Project 1



### Group Members:

- Aryan Adinath Popalghat (210020088)
  - Divyam Gupta (210020044)
  - Durgesh Sahane (210110096)

March 6, 2024

Indian Institute of Technology, Bombay

# 1. Introduction

In this project, we aim to develop a model for image classification on the CUB-200-2011 dataset. The CUB-200-2011 dataset is a widely used benchmark dataset containing images of 200 bird species, with a total of 11,788 images. Each image is annotated with the corresponding bird species label.

One of the key challenges in building deep learning models is balancing model complexity and performance. While deep neural networks are capable of achieving state-of-the-art results, they often come with a large number of parameters, leading to increased computational and memory requirements. To address this challenge, we impose a constraint on our model to have **less than 10 million parameters** while still achieving competitive performance on the classification task.

Our approach involves designing and training a convolutional neural network (CNN) architecture tailored to the characteristics of the CUB-200-2011 dataset. We explore various architectural choices, including the number of layers, filter sizes, and pooling strategies, to strike a balance between model complexity and performance. Additionally, we employ techniques such as regularization and optimization to enhance the generalization ability of our model and prevent overfitting.

The success of our project has implications beyond the specific task of bird classification. By demonstrating the feasibility of building a high-performing model with limited parameters, we contribute to the broader goal of developing efficient and scalable deep learning solutions for real-world applications.

In this report, we provide a detailed overview of our methodology, experimental setup, results, and analysis. We also discuss the implications of our findings and suggest directions for future research in the field of efficient deep learning models for image classification.

## 2. Data Preprocessing

The success of any machine learning model heavily relies on the quality and suitability of the dataset used for training. In this section, we outline the preprocessing steps applied to the CUB-200-2011 dataset to ensure its compatibility with our classification task and to enhance the performance and robustness of our model.

### 1. Dataset Overview:

The CUB-200-2011 dataset consists of a collection of bird images, each associated with a specific species label. Before proceeding with preprocessing, we performed an initial exploration of the dataset to understand its structure, distribution of classes, and image characteristics.

### 2. Data Augmentation:

To increase the diversity and variability of our training data and to prevent overfitting, we applied data augmentation techniques. These techniques include random rotation, horizontal flipping, and random cropping. Data augmentation helps the model learn invariant features and improves its generalization ability.

### 3. Image Rescaling and Normalization:

All images in the dataset were rescaled to a fixed size to ensure uniformity across the dataset and to facilitate efficient training. Additionally, we performed pixel normalization to standardize the pixel values of the images, typically by subtracting the mean and dividing by the standard deviation. Normalization helps stabilize the training process and accelerates convergence.

### 4. Data Splitting:

We partitioned the dataset into training, validation, and test sets. The training set is used to train the model, and the test set is used to evaluate the final model performance. We ensured that the distribution of classes is balanced across all sets to avoid biased model evaluation.

## 5. Data loading and Batching:

We utilized PyTorch's data loading utilities to efficiently load and preprocess the dataset. Data loaders were employed to generate mini-batches of data during training, allowing for parallel processing and memory optimization.

# 3. Model

We employed the ResNet-18 architecture as our baseline model and customized it to meet our project's constraints of less than 10 million parameters. The ResNet-18 architecture consists of a series of convolutional layers with residual connections, facilitating deeper networks while mitigating the vanishing gradient problem. To meet our parameter constraint, we reduced the number of filters and depth of the network, optimized regularization techniques, and carefully selected hyperparameters. This modified architecture retains the essence of ResNet-18 while ensuring computational efficiency and competitive performance on the CUB-200-2011 dataset.

## 1. Architecture:

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 112, 112]	4,704
BatchNorm2d-2	[-1, 32, 112, 112]	64
ReLU-3	[-1, 32, 112, 112]	0
MaxPool2d-4	[-1, 32, 56, 56]	0
Conv2d-5	[-1, 32, 56, 56]	9,216
BatchNorm2d-6	[-1, 32, 56, 56]	64
ReLU-7	[-1, 32, 56, 56]	0
Conv2d-8	[-1, 32, 56, 56]	9,216
BatchNorm2d-9	[-1, 32, 56, 56]	64
ReLU-10	[-1, 32, 56, 56]	0

ResidualBlock-11	[-1, 32, 56, 56]	0
Conv2d-12	[-1, 32, 56, 56]	9,216
BatchNorm2d-13	[-1, 32, 56, 56]	64
ReLU-14	[-1, 32, 56, 56]	0
Conv2d-15	[-1, 32, 56, 56]	9,216
BatchNorm2d-16	[-1, 32, 56, 56]	64
ReLU-17	[-1, 32, 56, 56]	0
ResidualBlock-18	[-1, 32, 56, 56]	0
Conv2d-19	[-1, 64, 28, 28]	18,432
BatchNorm2d-20	[-1, 64, 28, 28]	128
ReLU-21	[-1, 64, 28, 28]	0
Conv2d-22	[-1, 64, 28, 28]	36,864
BatchNorm2d-23	[-1, 64, 28, 28]	128
Conv2d-24	[-1, 64, 28, 28]	2,048
BatchNorm2d-25	[-1, 64, 28, 28]	128
ReLU-26	[-1, 64, 28, 28]	0
ResidualBlock-27	[-1, 64, 28, 28]	0
Conv2d-28	[-1, 64, 28, 28]	36,864
BatchNorm2d-29	[-1, 64, 28, 28]	128
ReLU-30	[-1, 64, 28, 28]	0
Conv2d-31	[-1, 64, 28, 28]	36,864
BatchNorm2d-32	[-1, 64, 28, 28]	128
ReLU-33	[-1, 64, 28, 28]	0
ResidualBlock-34	[-1, 64, 28, 28]	0
Conv2d-35	[-1, 128, 14, 14]	73,728
BatchNorm2d-36	[-1, 128, 14, 14]	256
ReLU-37	[-1, 128, 14, 14]	0
Conv2d-38	[-1, 128, 14, 14]	147,456
BatchNorm2d-39	[-1, 128, 14, 14]	256
Conv2d-40	[-1, 128, 14, 14]	8,192
BatchNorm2d-41	[-1, 128, 14, 14]	256
ReLU-42	[-1, 128, 14, 14]	0
ResidualBlock-43	[-1, 128, 14, 14]	0

Conv2d-44	[-1, 128, 14, 14]	147,456
BatchNorm2d-45	[-1, 128, 14, 14]	256
ReLU-46	[-1, 128, 14, 14]	0
Conv2d-47	[-1, 128, 14, 14]	147,456
BatchNorm2d-48	[-1, 128, 14, 14]	256
ReLU-49	[-1, 128, 14, 14]	0
ResidualBlock-50	[-1, 128, 14, 14]	0
Conv2d-51	[-1, 256, 7, 7]	294,912
BatchNorm2d-52	[-1, 256, 7, 7]	512
ReLU-53	[-1, 256, 7, 7]	0
Conv2d-54	[-1, 256, 7, 7]	589,824
BatchNorm2d-55	[-1, 256, 7, 7]	512
Conv2d-56	[-1, 256, 7, 7]	32,768
BatchNorm2d-57	[-1, 256, 7, 7]	512
ReLU-58	[-1, 256, 7, 7]	0
ResidualBlock-59	[-1, 256, 7, 7]	0
Conv2d-60	[-1, 256, 7, 7]	589,824
BatchNorm2d-61	[-1, 256, 7, 7]	512
ReLU-62	[-1, 256, 7, 7]	0
Conv2d-63	[-1, 256, 7, 7]	589,824
BatchNorm2d-64	[-1, 256, 7, 7]	512
ReLU-65	[-1, 256, 7, 7]	0
ResidualBlock-66	[-1, 256, 7, 7]	0
AdaptiveAvgPool2d-67	[-1, 256, 1, 1]	0
Linear-68	[-1, 200]	51,400

---



---

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 112, 112]	4,704
BatchNorm2d-2	[-1, 32, 112, 112]	64
ReLU-3	[-1, 32, 112, 112]	0
MaxPool2d-4	[-1, 32, 56, 56]	0
Conv2d-5	[-1, 32, 56, 56]	9,216
BatchNorm2d-6	[-1, 32, 56, 56]	64
ReLU-7	[-1, 32, 56, 56]	0
Conv2d-8	[-1, 32, 56, 56]	9,216
BatchNorm2d-9	[-1, 32, 56, 56]	64
ReLU-10	[-1, 32, 56, 56]	0
ResidualBlock-11	[-1, 32, 56, 56]	0
Conv2d-12	[-1, 32, 56, 56]	9,216
BatchNorm2d-13	[-1, 32, 56, 56]	64
ReLU-14	[-1, 32, 56, 56]	0
Conv2d-15	[-1, 32, 56, 56]	9,216
BatchNorm2d-16	[-1, 32, 56, 56]	64
ReLU-17	[-1, 32, 56, 56]	0
ResidualBlock-18	[-1, 32, 56, 56]	0
Conv2d-19	[-1, 64, 28, 28]	18,432
BatchNorm2d-20	[-1, 64, 28, 28]	128
ReLU-21	[-1, 64, 28, 28]	0
Conv2d-22	[-1, 64, 28, 28]	36,864
BatchNorm2d-23	[-1, 64, 28, 28]	128
Conv2d-24	[-1, 64, 28, 28]	2,048
BatchNorm2d-25	[-1, 64, 28, 28]	128
ReLU-26	[-1, 64, 28, 28]	0
ResidualBlock-27	[-1, 64, 28, 28]	0
Conv2d-28	[-1, 64, 28, 28]	36,864
BatchNorm2d-29	[-1, 64, 28, 28]	128
ReLU-30	[-1, 64, 28, 28]	0
Conv2d-31	[-1, 64, 28, 28]	36,864
BatchNorm2d-32	[-1, 64, 28, 28]	128
ReLU-33	[-1, 64, 28, 28]	0
ResidualBlock-34	[-1, 64, 28, 28]	0
Conv2d-35	[-1, 128, 14, 14]	73,728
BatchNorm2d-36	[-1, 128, 14, 14]	256
ReLU-37	[-1, 128, 14, 14]	0

Conv2d-38	[-1, 128, 14, 14]	147,456
BatchNorm2d-39	[-1, 128, 14, 14]	256
Conv2d-40	[-1, 128, 14, 14]	8,192
BatchNorm2d-41	[-1, 128, 14, 14]	256
ReLU-42	[-1, 128, 14, 14]	0
ResidualBlock-43	[-1, 128, 14, 14]	0
Conv2d-44	[-1, 128, 14, 14]	147,456
BatchNorm2d-45	[-1, 128, 14, 14]	256
ReLU-46	[-1, 128, 14, 14]	0
Conv2d-47	[-1, 128, 14, 14]	147,456
BatchNorm2d-48	[-1, 128, 14, 14]	256
ReLU-49	[-1, 128, 14, 14]	0
ResidualBlock-50	[-1, 128, 14, 14]	0
Conv2d-51	[-1, 256, 7, 7]	294,912
BatchNorm2d-52	[-1, 256, 7, 7]	512
ReLU-53	[-1, 256, 7, 7]	0
Conv2d-54	[-1, 256, 7, 7]	589,824
BatchNorm2d-55	[-1, 256, 7, 7]	512
Conv2d-56	[-1, 256, 7, 7]	32,768
BatchNorm2d-57	[-1, 256, 7, 7]	512
ReLU-58	[-1, 256, 7, 7]	0
ResidualBlock-59	[-1, 256, 7, 7]	0
Conv2d-60	[-1, 256, 7, 7]	589,824
BatchNorm2d-61	[-1, 256, 7, 7]	512
ReLU-62	[-1, 256, 7, 7]	0
Conv2d-63	[-1, 256, 7, 7]	589,824
BatchNorm2d-64	[-1, 256, 7, 7]	512
ReLU-65	[-1, 256, 7, 7]	0
ResidualBlock-66	[-1, 256, 7, 7]	0
AdaptiveAvgPool2d-67	[-1, 256, 1, 1]	0
Linear-68	[-1, 200]	51,400
=====		

## 2. Parameters:

---

---

Total params: 2,850,280

Trainable params: 2,850,280

Non-trainable params: 0

---

Input size (MB): 0.57

Forward/backward pass size (MB): 31.39



Params size (MB): 10.87  
Estimated Total Size (MB): 42.84

```
=====
Total params: 2,850,280
Trainable params: 2,850,280
Non-trainable params: 0
-----

Input size (MB): 0.57
Forward/backward pass size (MB): 31.39
Params size (MB): 10.87
Estimated Total Size (MB): 42.84
-----
```


## 4. Training


- 1. Hyperparameters:
  - Number of classes: 200
  - Batch size: 32
  - Number of epochs: 30

- 2. Training Cost and Accuracy:  
Performance Summary

Epoch	Training Loss	Training Accuracy
1	0.16568245897195996	0.8908371040723982
2	0.15619492092553308	2.0786199095022626
3	0.14885187526633836	3.0118778280542986
4	0.1415379787057773	4.3127828054298645
5	0.1339182592422714	5.896493212669683
6	0.12621298476177104	8.79524886877828
7	0.1168067765114534	12.471719457013574
8	0.10929018815313529	16.43099547511312
9	0.10088828418454433	20.9841628959276
10	0.09257651808170173	24.90101809954751
11	0.08631974596917899	29.920814479638008
12	0.07737295437452481	35.39309954751131

13	0.06975575612343814	40.540158371040725
14	0.06229199597199039	47.171945701357465
15	0.054177477770396484	52.37556561085973
16	0.04498270105958255	60.39309954751131
17	0.03572620590516615	68.35407239819004
18	0.027278319380469453	76.96549773755656
19	0.018614862843235425	85.98699095022624
20	0.013246690892357362	90.2997737556561
21	0.008470857512770533	94.61255656108597
22	0.0051170337799835396	97.69513574660634
23	0.0027559062232141176	99.27884615384616
24	0.0016490024621429857	99.6747737556561
25	0.0012457637683640137	99.77375565610859
26	0.0009343491161844748	99.90101809954751
27	0.0008055211423592836	99.84445701357465
28	0.010445750839702735	90.69570135746606
29	0.01878154013154194	81.53280542986425
30	0.004308475319473115	96.74773755656109





## 5. Results

1. **Training Accuracy:** Accuracy of the network on the train images: 99.43438914027149%
2. **Test Accuracy:** Accuracy of the network on the test images: 26.95080576759966%
3. Got 7066 correct images out of 7072 total images (99.92% accuracy) (For training data)
4. Got 1351 correct images out of 4716 total images (28.65% accuracy) (For testing data)

Note: The NewModel12.ipynb submitted in the folder is before addition of checkpoints in it. Because checkpoint calculation is later done on a different device.